



Introduction to NoSQL

COMP23111 – Database Systems

Gareth Henshall

Department of Computer Science

Introduction to NoSQL Systems

Who Uses Not Only SQL (NOSQL) Systems



amazon



Introduction to NoSQL Systems

Emergence of NoSQL Systems

- ✍ The need by organisations to **store vast** amount of data, for example, emails, tweets, posts, updates, pictures, etc.
- ✍ SQL **service overload**, such as **powerful query language, concurrency control**, etc. which most of these applications do not need.
- ✍ Hence, Some of these **organisations** developed their own systems referred to as **NoSQL** systems to effectively manage these **vast** amount of data.

Introduction to NoSQL Systems

Emergence of NoSQL Systems (Cont.)

Google BigTable



Column-base



amazon
DynamoDB

Key-value and column based

Facebook



cassandra

Key-value based



CouchDB
relax



mongoDB®

Document-based



GraphBase



neo4j

Graph-based

Introduction to NoSQL Systems

Characteristics Of NoSQL Systems

✍ NoSQL Characteristics Related to Distributed Database and Distributed Systems

- ✍ Scalability (horizontal scalability)
- ✍ Availability, Replication and Eventual Consistency
- ✍ Replication Model (master-slave replication & master-master replication)
- ✍ Sharding of files
- ✍ High-Performance Data Access


✍ NoSQL Characteristics Related to Data Models and Query Language

- ✍ Not Requiring a Schema (Schemaless)
- ✍ Less Powerful Query Language
- ✍ Versioning

Introduction to NOSQL Systems

Category of NoSQL Systems


Key-value based NoSQL



Example: Riak, Redis server, Scalaris

Possess a simple data model based on fast access key to value associated with the key. The value can be record or object or document or complex data structure. Works best for shopping cart contents, etc.


Document-based NoSQL



Example: MongoDB, CouchDB, OrientDB

Stores data in form of documents using well known data format, such as JSON(JavaScript object Notation) accessible via unique document "id". Mostly used for CMS, blogging, eCommerce, etc.


Column-based NoSQL



Example: BigTable, Cassandra, Hbase

Partition table by column into column family, where each column is stored in its own file. Widely used for managing large data storage, such as data warehouse, etc.

Graph-based NoSQL



Example: Neo4J, InfoGrid, FlockDB

Stores entities as nodes and relation between entities as edges, where every node and edge has a unique identifier and can be traversed using path expression. Mostly used for social network, etc.

The CAP Theorem

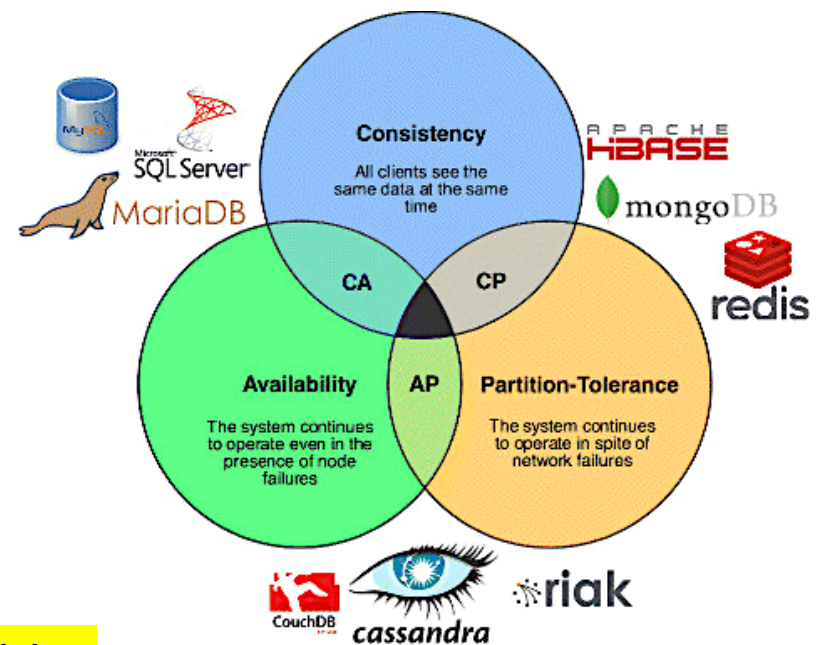
Consistency, Availability and Partition Tolerance

✍ **Consistency** implies all nodes should contain same copies of replicated data item visible for various transactions.

✍ **Availability** implies that the system should be consistently available for read and write operations.

✍ **Partition tolerance** implies that the system should be continually available in situations where the system is partitioned by network fault

Note: The cap theorem states that it is not possible to guarantee all the desirable properties – consistency, availability and partition tolerance



Ref []

The CAP Theorem

Eventual Consistency

- ✍ In a NoSQL distributed database, a **weaker** consistency is often acceptable, while guaranteeing **availability** and **partition** tolerance - where the database will **eventually be consistent**.
- ✍ Hence the **BASE** property is used instead of the **ACID** property in centralised database.
- ✍ BASE: **B**asically **A**vailable, **S**oft state, **E**ventual consistency
 - ✍ **B**asically **A**vailable means the DDB is available following the CAP theorem
 - ✍ **S**oft state means the system's state may change even transaction execution
 - ✍ **E**ventual constancy implies that the system will become consistent over time

Document-Based NoSQL Systems: MongoDB



- ✍ **MongoDB** is a **general purpose, document-based, distributed database** built for **modern application** developers and for the **cloud** era.
- ✍ **MongoDB** is developed by **MongoDB Inc** and licensed under the Server Side Public License (SSPL).


DB Systems Rankings

358 systems in ranking, September 2020


Rank			DBMS	Database Model	Score		
Sep 2020	Aug 2020	Sep 2019			Sep 2020	Aug 2020	Sep 2019
1.	1.	1.	Oracle +	Relational, Multi-model i	1369.36	+14.21	+22.71
2.	2.	2.	MySQL +	Relational, Multi-model i	1264.25	+2.67	-14.83
3.	3.	3.	Microsoft SQL Server +	Relational, Multi-model i	1062.76	-13.12	-22.30
4.	4.	4.	PostgreSQL +	Relational, Multi-model i	542.29	+5.52	+60.04
5.	5.	5.	MongoDB +	Document, Multi-model i	446.48	+2.92	+36.42
6.	6.	6.	IBM Db2 +	Relational, Multi-model i	161.24	-1.21	-10.32
7.	7.	↑ 8.	Redis +	Key-value, Multi-model i	151.86	-1.02	+9.95
8.	8.	↓ 7.	Elasticsearch +	Search engine, Multi-model i	150.50	-1.82	+1.23
9.	9.	↑ 11.	SQLite +	Relational	126.68	-0.14	+3.31
10.	↑ 11.	10.	Cassandra +	Wide column	119.18	-0.66	-4.22

<https://db-engines.com/en/ranking>

DB Rankings 2013 - 2020



DB Rankings 2013 - 2020



- ✍ **Popularity**
- ✍ **High performance**
- ✍ **High availability**
- ✍ **Horizontal Scalability**
- ✍ **Rich query language**
- ✍ **Support for Multiple Storage Engines**

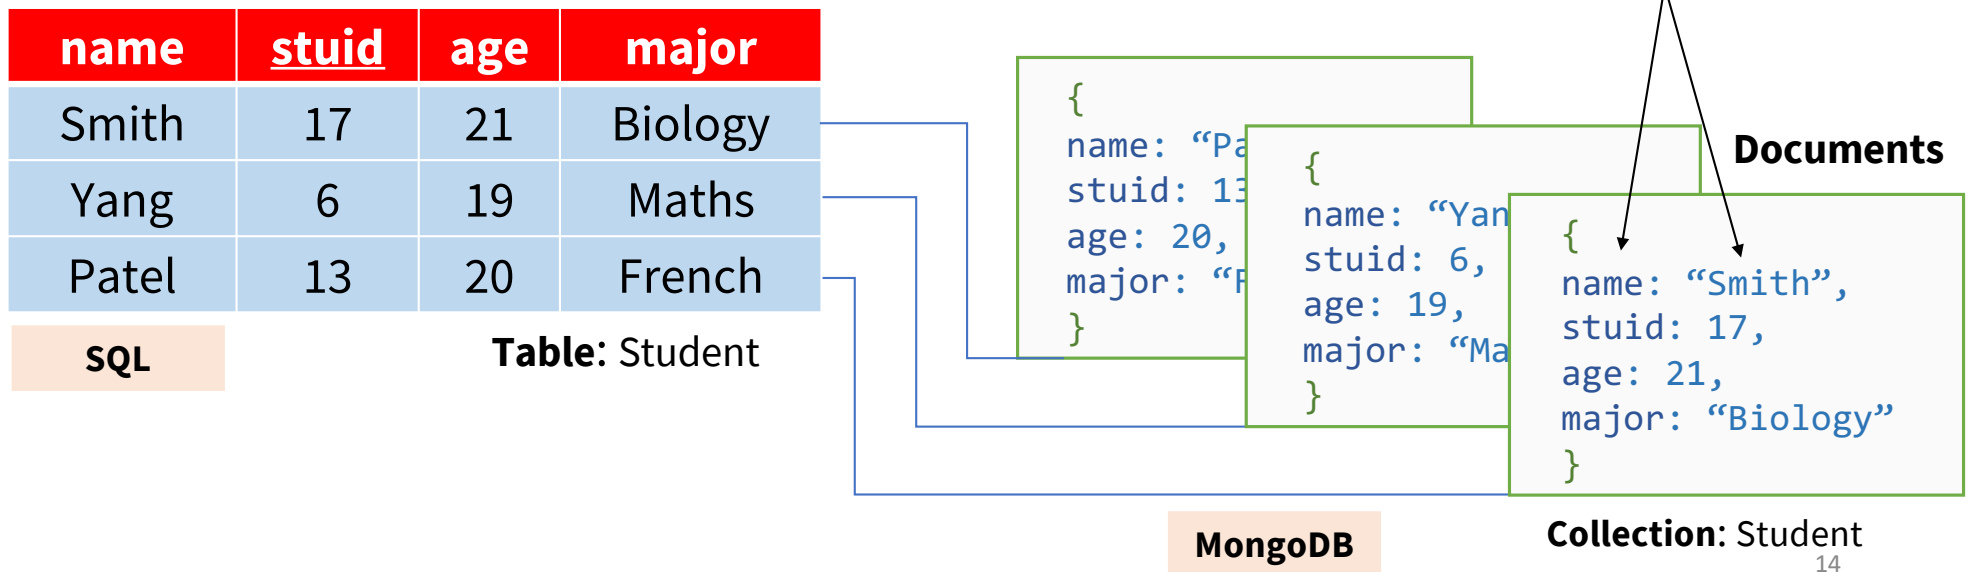
Document-Based NoSQL Systems: MongoDB

Terminology and Concepts: SQL and MongoDB

SQL Terms/ Concepts	MongoDB Terms/ Concepts
database	database
table	collection
row	document
column	field
index	index
table joins	\$lookup
primary key	primary key

MongoDB Database and Collection

In MongoDB, **database** holds **collection** of documents which are in **JSON-style** (JavaScript Object Notation Style) format.



MongoDB Database and Collection

Explicitly create a Database

```
db.createCollection(<name>, <options>)  
  
db.createCollection("myNewDB")
```

The `myNewDB` explicitly creates a new database if it does not already exist.

Create a Collection

```
db.myNewCollection1.insertOne({name: "Patel"})
```

The `myNewCollection` creates a new collection if it does not already exist.

Create Database and Collection

```
use myNewDB  
  
db.myNewCollection2.insertOne({name: "Smith"})
```

The `insertOne()` operation creates both the database `myNewDB` and the collection `myNewCollection` if they do not already exist.

MongoDB Document

BSON and JSON

✍ In MongoDB, **database** holds **collection** of documents which are in **JSON-style** (JavaScript Object Notation Style) format.

✍ **JSON** is an open, **human** and **machine-readable** standard to transmit data objects consisting of **attributes-value pairs**.

✍ MongoDB represents JSON documents in **binary-encoded** format called **Binary JSON** (BSON - *[bee · sahn]*) behind the scene.

MongoDB Document

JavaScript Object Notation Structure: Example

✍ Easy for **humans** to write and read,
and easy for **computers** to parse and
generate

✍ **Objects** can be nested

✍ Built on

✍ name/value pairs

✍ Ordered list of values

```
{  
  "Title": "The Cuckoo's Calling"  
  "Author": "Robert Galbraith",  
  "Genre": "classic crime novel",  
  "Detail": {  
    "Publisher": "Little Brown"  
    "Publication_Year": 2013,  
    "ISBN-13": 9781408704004,  
    "Language": "English",  
    "Pages": 494  
  }  
  "Price": [  
    {  
      "type": "Hardcover",  
      "price": 16.65,  
    }  
    {  
      "type": "Kindle Edition",  
      "price": 7.03,  
    }  
  ]  
}
```

Diagram annotations:

- Object Starts (at the opening curly brace)
- Object Starts (at the opening curly brace of the nested object)
- Value string (pointing to "Little Brown")
- Value number (pointing to 2013)
- Object ends (at the closing curly brace of the nested object)
- Array starts (at the opening square bracket)
- Object Starts (at the opening curly brace of the first array element)
- Object ends (at the closing curly brace of the first array element)
- Object Starts (at the opening curly brace of the second array element)
- Object ends (at the closing curly brace of the second array element)
- Array ends (at the closing square bracket)
- Object ends (at the closing curly brace)

MongoDB Document

Binary JSON (BSON)

✍ **Binary-encoded** serialization of
JSON-like docs

✍ Goals:

✍ **Lightweight** (Keeping spatial overhead to a minimum)

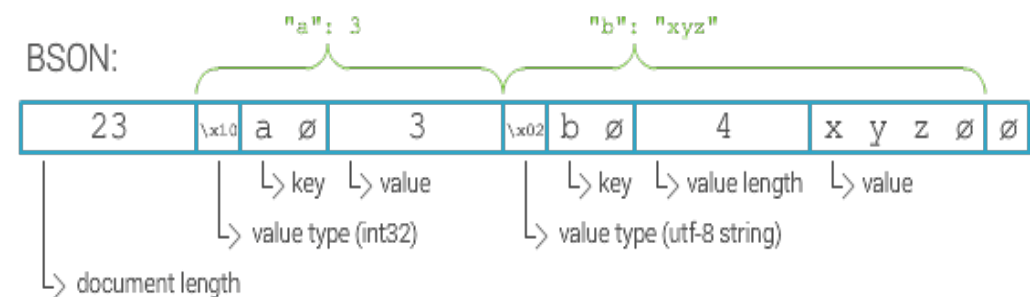
✍ **Traversable** (stores the length of values to find that specific key)

✍ **Efficient** (decoding and encoding)

JSON:

```
{  
  "a": 3,  
  "b": "xyz"  
}
```

BSON:



MongoDB Document

Binary JSON (BSON) (Cont.)

In BSON:

✍ Data is represented in **field / value** pairs

✍ A field/value pair consists of a field name followed by a colon, followed by a value:

Example: `name: "Joseph"`

✍ Fields are separated by commas

Example: `name: "Joseph", "course": "Databases", score: 80`

✍ Curly braces hold objects (documents)

Example: `{name: "Joseph", "course": "Databases", score: 80}`

MongoDB Document

Binary JSON (BSON) (Cont.)

✍ **Embedded** document

Example: `{name: "Joseph", courses: {course1: "databases", course2: "programming"}}`

The field **courses** holds an **embedded document** containing the fields **course1** and **course2**

✍ An **array** is stored in brackets []

Example: `{name: "Joseph", courses: ["databases", "programming"]}`

✍ An **array** of embedded document contains documents embedded in the array

Example: `{name: "Joseph", courses: [{course1: "databases", course2: "programming"}]}`

MongoDB Document

Document Structure: Example

```
db.projectCollection.insertOne(  
{  
  pname: "ProductX",  
  plocation: "Frankfurt",  
  staff: [  
    {fname: "John", lname: "Smith", hours: 28.4},  
    {fname: "Joyce", lname: "Marry", hours: 23.4}  
  ],  
  supervisors: ["James Brown", "Louis Lampard"],  
  status: {finished: 1, ongoing: 0, comment: "none"}  
})
```

Holds the objectId

```
_id: ObjectId("5dac5a50d16c90c5b6007ac8")  
pname: "ProductX"  
plocation: "Frankfurt"  
pyears: 3  
✓ staff: Array  
  ✓ 0: Object  
    fname: "John"  
    lname: "Smith"  
    hours: 28.4  
  ✓ 1: Object  
    fname: "Joyce"  
    lname: "Marry"  
    hours: 23.4  
✓ supervisors: Array  
  0: "James Brown"  
  1: "Louis Lampard"  
✓ status: Object  
  finished: 1  
  ongoing: 0  
  comment: "none"
```

MongoDB Document

Document Structure: Example

```
db.projectCollection.insertOne(  
{  
  pname: "ProductX",  
  plocation: "Frankfurt",  
  pyears: 3  
  staff: [  
    {fname: "John", lname: "Smith", hours: 28.4},  
    {fname: "Joyce", lname: "Marry", hours: 23.4}  
  ],  
  supervisors: ["James Brown", "Louis Lampard"],  
  status: {finished: 1, ongoing: 0, comment: "none"}  
})
```

```
{ "_id" : ObjectId("5dabb758d16c90c5b6007ac4"),  
  "pname": "ProductX",  
  "plocation" : "Frankfurt",  
  "pyears": 3  
  "staff" : [  
    { "fname" : "John", "lname" : "Smith", "hours" : 28.4 },  
    { "fname" : "Joyce", "lname" : "Marry", "hours" : 23.4 } ],  
  "supervisors" : [ "James Brown", "Louis Lampard" ],  
  "status" : { "finished" : 1, "ongoing" : 0, "comment" :  
    "none" } }
```

BSON

MongoDB Document

The `_id` Field

- ✍ Each document in a collection requires a unique `_id` field
- ✍ The `_id` field acts as a primary key to the collection
- ✍ If an inserted document omits the `_id` field, an **Objectid** is automatically generated for the `_id` field
- ✍ The `_id` field is always the first field in the document
- ✍ The `_id` field may contain BSON data type except arrays

MongoDB CRUD Operations

Query Operators

Name	Description
<code>\$eq</code>	Matches values that are equal to a specified value.
<code>\$gt</code>	Matches values that are greater than a specified value.
<code>\$gte</code>	Matches values that are greater than or equal to a specified value.
<code>\$in</code>	Matches any of the values specified in an array.
<code>\$lt</code>	Matches values that are less than a specified value.
<code>\$lte</code>	Matches values that are less than or equal to a specified value.
<code>\$ne</code>	Matches all values that are not equal to a specified value.
<code>\$nin</code>	Matches none of the values specified in an array.

MongoDB CRUD Operations

Create Operations

✍ Create or insert operations add **new document** to a **collection**

✍ If the collection does not currently **exist**, the insert operations **create** it

✍ The following **methods** are used to insert documents into a collection:

```
db.collectionName.insertOne(<document>)
```

Insert a **single** document into a collection

```
db.collectionName.insertMany(<documents>)
```

Insert a **multiple** document into a collection

```
db.collectionName.insert(<documents>)
```

Insert a **single** or **multiple** document into a collection

MongoDB CRUD Operations

Read Operations

✍ Read operation retrieves documents from a collection – queries a collection for documents

✍ The following **methods** are used to read document in a collection

```
db.collectionName.find(<query> <projection>)
```

query specifies selection criteria using query operators, and **projection** specifies the **field** to return in the **document**

```
db.projectCollection.find( { "pyears": { $gt:3 } }, { staff:1, supervisors:1 })
```

Displays all **staff** and **supervisors** who **work** on or **supervise** projects that are **over three years**

```
db.projectCollection.find( {"staff.hours": { $gt:18 }} ,{staff:1} )
```

Displays all **staff** who has worked for over **18** hours on any **project**

MongoDB CRUD Operations

Update Operations

✍️ Modify existing document in a collection

✍️ The following methods are used to update documents in a collection

```
db.collectionName.updateMany (<filter>, <update action>, <options>)
```

filter denotes the selection criteria for the update, **update action** implies the modification to apply, and **options** implies additional optional actions that can be assigned to the operation .

```
db.projectCollection.updateMany(  
  { "pyears": { $gte: 5 } },  
  {$set: { "status.finished": 1, "status.ongoing": 0}}  
)
```

Updates the **finished** and **ongoing** status of all **projects** whose number of **years** is **greater** than or **equal** to **five** to **1** and **0** respectively

MongoDB CRUD Operations

Delete Operations

✍ Delete operations remove documents from a collection.

✍ The following methods are used to remove documents from a collection

```
db.collectionName.deleteMany (<filter>, <options>)
```

filter specifies deletion criteria using query operators, and **options** specifies additional optional actions that can be assigned to the operation.

```
db.projectCollection.deleteMany(  
    { "status.finished" :1, "status.ongoing": 0}  
)
```

Deletes all **projects** whose **finished** status is **1** and **ongoing** status is **0**

MongoDB CRUD Operations

More Operations Example

```
SELECT staff, supervisors  
FROM projectCollection  
WHERE pyears >=3 OR pyears <= 5
```

```
SELECT * FROM projectCollection  
WHERE pyears >=5 AND  
finished = 1 AND  
Ongoing = 0
```

```
db.projectCollection.find(  
  { $or: [{"pyears": { $gte:3 }},  
    {"pyears": { $lte:5 }} ]},  
  { staff:1, supervisors:1 })
```

```
db.projectCollection.find(  
  { $and: [{"pyears": { $gte:5 }},  
    {"status.finished": { $eq:1 }},  
    {"status.ongoing": { $eq:0 }} ]})
```

MongoDB CRUD Operations

Operations Output Sample 1

```
db.projectCollection.find(  
  { $or: [{"pyears": { $gte:3 }},  
    {"pyears": { $lte:5 }} ]},  
  { staff:1, supervisors:1 })
```



```
{  
  "_id":ObjectId(  "5dac5a50d16c90c5b6007ac8"  ),  
  "staff":[  
    {  
      "fname":"John",  
      "lname":"Smith",  
      "hours":28.4  
    },  
    {  
      "fname":"Joyce",  
      "lname":"Marry",  
      "hours":23.4  
    }  
  ],  
  "supervisors":[  
    "James Brown",  
    "Louis Lampard"  
  ]  
}
```

MongoDB CRUD Operations

Operations Output Sample 2

```
db.projectCollection.find(  
  { $and: [{"pyears": { $gte:5 }},  
    {"status.finished": { $eq:1 }},  
    {"status.ongoing": { $eq:0 }} ]})
```



```
{  
  "_id":ObjectId("5dac5bdcd16c90c5b6007aca"),  
  "pname":"ProductY",  
  "plocation":"London",  
  "pyears":5,  
  "staff":[  
    {  
      "fname":"Salah",  
      "lname":"Hammed",  
      "hours":89.4  
    },  
    {  
      "fname":"Marry",  
      "lname":"Fin",  
      "hours":98.4  
    }  
  ],  
  "supervisors":[  
    "James Brown",  
    "Louis Lampard"  
  ],  
  "status":{  
    "finished":1,  
    "ongoing":0,  
    "comment":"none",  
    "finish":1  
  }  
}
```