



Transaction Processing

Schedules, Conflicts and Serializability

Dr Stewart Blakeway

Lecturer in Computer Science

Intro

Learning Objectives

- Desirable Properties of Transactions (ACID)
- Schedules
 - Conflicting Operations in a Schedule
 - Serializable Schedules
 - Conflict Equivalence of Two Schedules
- Testing the Serializability of a Schedule

Desirable Properties of Transactions

- Transactions should possess several properties which are called the ACID properties
- Atomicity
- Consistency Preservation
- Isolation
- Durability or Permanency

Schedules

- When executing transactions concurrently in an interleaved fashion we need to create a schedule
- A schedule is defined as: S of n transaction T_1, T_2, \dots, T_n
- A shorthand notation describing a schedule uses symbols b, r, w, e, c and a

Example

$S_a: r_1(X); r_2(X); w_1(X); r_1(Y); w_2(X); w_1(Y);$

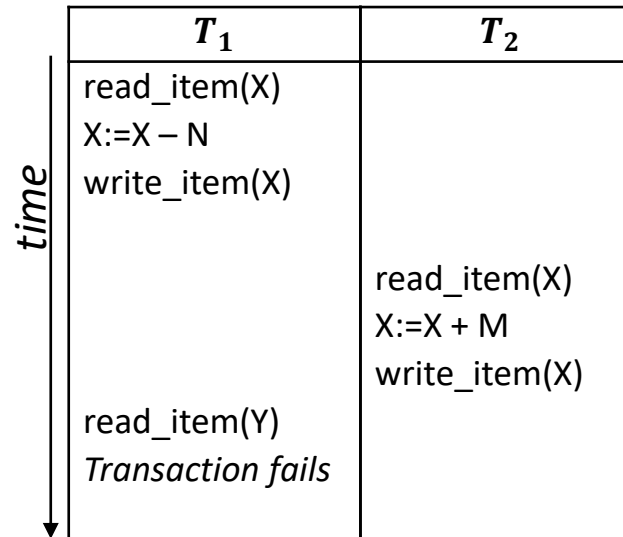
	T_1	T_2
time ↓	read_item(X) $X := X - N$	read_item(X) $X := X + M$
	write_item(X) read_item(Y)	write_item(X)
	$Y := Y + N$ write_item(Y)	

$S_b: r_1(X); w_1(X); r_2(X); w_2(X); r_1(Y); a_1;$

Your Turn

$S_b:$

w, r, c and a



Conflicting Operations in a Schedule

- Two operations in a schedule are said to conflict if they satisfy all three conditions
 - Condition 1: They belong to different transactions
 - Condition 2: They access the same item X
 - Condition 3: At least one of the operations is a write
- Conflicting Operations in a Schedule could result in inconsistent data!

Example

- 1: They belong to different transactions
- 2: They access the same item X
- 3: At least one of the operations is a write

read-write conflict

$S_a: r_1(X); r_2(X); w_1(X); r_1(Y); w_2(X); w_1(Y);$

write-write conflict

time	T_1	T_2
	read_item(X) $X := X - N$ write_item(X) read_item(Y) $Y := Y + N$ write_item(Y)	read_item(X) $X := X + M$ write_item(X)

A Complete Schedule

- The goal is to establish a complete schedule, which is defined as:
 1. The operations in S are exactly those operations in T_1, T_2, \dots, T_n
 2. For any pair of operations from the same transaction, their relative order in S is preserved
 3. For any two conflicting operations, one of the two must occur before the other in the schedule

The easy approach

<i>time</i> ↓	T_1	T_2
	read_item(X)	
	$X := X - N$	
	write_item(X)	
	read_item(Y)	
	$Y := Y + N$	
	write_item(Y)	
		read_item(X)
		$X := X + M$
		write_item(X)

Schedule A

<i>time</i> ↓	T_1	T_2
		read_item(X)
		$X := X + M$
		write_item(X)
	read_item(X)	
	$X := X - N$	
	write_item(X)	
	read_item(Y)	
	$Y := Y + N$	
	write_item(Y)	

Schedule B

These are examples of Serial Schedules

One runs after the other in series.

With two transactions there can only be two possible schedules using a serial schedule.

Unacceptable in practice.

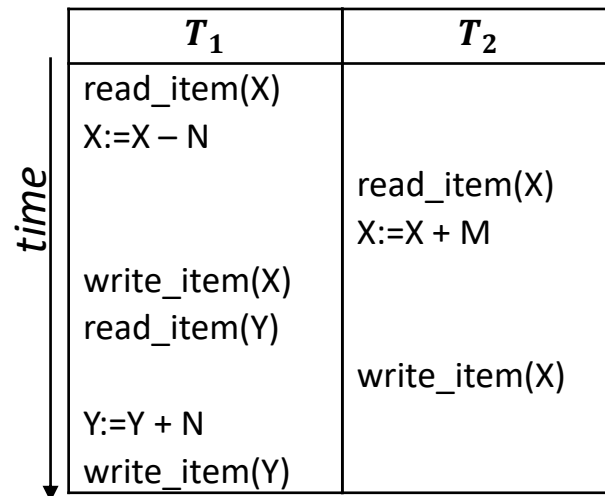
Long I/O waits (no switching)
A large transaction might hog the CPU

Allowing Interleaving


These are examples of
NonSerial Schedules

They are interleaved and run
concurrently.

- Ideally we want to allow interleaving where possible
- Interleaving of Transactions results in many possible schedules, here are two examples.



Schedule C



Schedule D

Quickly help me run through the schedules
with values: $X=90$, $Y=90$, $N=3$, $M=2$.

Correct or not correct

	T_1	T_2
time ↓	read_item(X)	
	$X := X - N$	
	write_item(X)	
	read_item(Y)	
	$Y := Y + N$	
	write_item(Y)	
		read_item(X)
		$X := X + M$
		write_item(X)

Schedule A

	T_1	T_2
time ↓		read_item(X)
		$X := X + M$
		write_item(X)
	read_item(X)	
	$X := X - N$	
	write_item(X)	
	read_item(Y)	
	$Y := Y + N$	
	write_item(Y)	

Schedule B

	T_1	T_2
time ↓	read_item(X)	
	$X := X - N$	
		read_item(X)
		$X := X + M$
	write_item(X)	
	read_item(Y)	
		write_item(X)
	$Y := Y + N$	
	write_item(Y)	

Schedule C

So...

- Serial Transactions always give us the correct result
- Non Serial is a bit hit or miss
 - But we want non serial because of concurrency
- What if there was a way to determine if a non serial schedule gives us the correct result, then our problems would be solved
 - This brings us to Serializable Schedules

Serializable Schedule

- A schedule S of n transactions is serializable if it is **equivalent** to some serial schedule of the same n transactions
- There are $n!$ possible schedules of serial transactions
 - Many more for non serial schedules
- But what do we mean by equivalent?
 - **Result Equivalent:** In the simplest terms, compare the results of a serial schedule to that of a non serial schedule.
 - But... two different schedules may accidentally produce the same final state

S_1
read_item(X)
$X := X + 10$
write_item(X)

S_2
read_item(X)
$X := X * 1.1$
write_item(X)

If $X=100$ they produce the same result. So are they equivalent? What if $X=20$?

Safe and General Approach to Equivalence

- The safe approach is to focus only on the read and write operations of the transactions.
 - Don't make assumptions about the other internal operations included within the transactions.
- Two schedules are equivalent if the operations applied to each data item affected by the schedules should be applied to that item in both schedules **in the same order**
- The most common approach to equivalence of schedules is the **conflict equivalence**

Conflict Equivalence of Two Schedules

- Two schedules are said to be conflict equivalent if the relative order of any two conflicting operations is the same in both schedules.
 - Different transactions
 - Access the same data item
 - Both or one is a write operation
- Not Conflicting Equivalence is when conflicting operations are applied in different orders in two schedules

Serializable Schedules

- S is serializable if its (conflict) equivalent to some serial schedule
- Therefore, by reordering the operations in a schedule it could become serializable

Schedule C (Non Serial)

	T_1	T_2
time ↓	read_item(X)	
	X:=X - N	
		read_item(X)
		X:=X + M
	write_item(X)	
	read_item(Y)	
	Y:=Y + N	
	write_item(X)	
	write_item(Y)	

Schedule A (Serial)

	T_1	T_2
time ↓	read_item(X)	
	X:=X - N	
	write_item(X)	
	read_item(Y)	
	Y:=Y + N	
	write_item(Y)	
		read_item(X)
		X:=X + M
		write_item(X)

Schedule D (Non Serial)




	T_1	T_2
time ↓	read_item(X)	
	X:=X - N	
	write_item(X)	
		read_item(X)
		X:=X + M
		write_item(X)
	read_item(Y)	
	Y:=Y + N	
	write_item(Y)	

Testing the Serializability of a Schedule

- We can use graph theory to easily test if a schedule is serializable or not
- Recall that we only focus on the read and write operations

Schedule C (Non Serial)

T_1	T_2
write_item(X) read_item(Y) write_item(Y)	read_item(X) write_item(X)





These two transactions conflict, we have a read operation and a write operation

Testing the Serializability of a Schedule

Schedule C (Non Serial)



T_1	T_2
write_item(X)	read_item(X)
read_item(Y)	
write_item(Y)	write_item(X)



Testing the Serializability of a Schedule

Schedule C (Non Serial)

T_1	T_2
write_item(X) read_item(Y) write_item(Y)	read_item(X) write_item(X)




These two transactions conflict, we have a write operation and a write operation

Because we have cycle, the schedule is not serializable

Testing the Serializability of a Schedule

Schedule D (Non Serial)

	T_1	T_2
time ↓	read_item(X)	
	write_item(X)	
		read_item(X)
		write_item(X)
	read_item(Y)	
	write_item(Y)	




We can now skip T_1
because we have
already found a conflict

Testing the Serializability of a Schedule

Schedule D (Non Serial)

	T_1	T_2
time ↓	read_item(X) write_item(X)	read_item(X) write_item(X)
	read_item(Y) write_item(Y)	



No conflict from T_2 to T_1 so this schedule is serializable

Let look at a more complicated example

Schedule


	T_1	T_2	T_3
time ↓	read_item(B)		read_item(C)
	read_item(A)		
	write_item(A)	write_item(A)	
	write_item(B)	write_item(B)	
			write_item(A)
			write_item(B)
			write_item(C)

Using the approach we have seen. Determine if this schedule is serializable or not

Let look at a more complicated example

Schedule


	T_1	T_2	T_3
time ↓	read_item(B)		
	read_item(A)		read_item(C)
	write_item(A)	write_item(A)	
		write_item(B)	
	write_item(B)		write_item(A)
			write_item(B)
			write_item(C)



Let look at a more complicated example

Schedule

	T_1	T_2	T_3
time ↓	read_item(B)		read_item(C)
	read_item(A)		
	write_item(A)	write_item(A)	
		write_item(B)	
	write_item(B)		write_item(A)
			write_item(B)
			write_item(C)




All good, no conflicts

Let look at a more complicated example

Schedule

	T_1	T_2	T_3
time ↓	read_item(B)		read_item(C)
	read_item(A)	write_item(A)	
	write_item(A)	write_item(B)	write_item(A)
	write_item(B)		write_item(B)
			write_item(C)



Cycle create between T_2 and T_1 , T_2 and T_3

Conclusion

- Desirable Properties of Transactions (ACID)
- Schedules
 - Conflicting Operations in a Schedule
 - Serializable Schedules
 - Conflict Equivalence of Two Schedules
- Testing the Serializability of a Schedule