

COMP24112: Machine Learning

Chapter 5: Loss Functions I

Dr. Tingting Mu

Email: tingting.mu@manchester.ac.uk



Content

- Typical approaches of constructing losses for **regression**
 - Non-probabilistic losses
 - Probabilistic losses



“Data + Model” in Supervised Learning

- “Data + Model” Strategy:



“Data + Model” in Supervised Learning

- After the learning (or called training) is finished, the final product is:

A trained model:

$$f(\theta(D_{tr}), x)$$

- The process of using the trained model on unseen data (or called query data, test data) is called inference.

$$answer = f(\theta(D_{tr}), x_{query})$$

Loss Function

- Loss function is essential in training, computed using the training data.
- It decides how good the model parameters are, how well the model fits your training data.
- Other names: error function, cost function, or more general, objective function.
- To train a machine learning model, you pick a loss function $O(\theta)$. Then:
 - Minimise it, if O evaluates how bad the model is.
 - Maximise it, if O evaluates how good the model is.
- Supervised learning: $O(\theta, D_{tr})$



Losses for Training Regression Models

Training Data: $D_{tr} = \{\mathbf{x}_i, y_i\}_{i=1}^N$

Feature vector: $\mathbf{x}_i \in R^d$

Target output: $\mathbf{y}_i \in R^c$ where $\mathbf{y}_i = [y_{i1}, y_{i2}, \dots, y_{ic}]$

Non-probabilistic Regression Losses

Regression Error based Loss

- Recall RMSE?
- Some regression losses are simplified versions of RMSE computed using training samples.

The prediction for each training sample is computed by $\hat{y}_i = f(\theta, \mathbf{x}_i)$.


– **Sum-of-squares error loss:**
$$O(\theta) = \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^c (\hat{y}_{ij} - y_{ij})^2$$

– **Mean-squared error loss:**

A single-output example ($c=1$):
$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$$

A Regression Example

- Fit a **linear model** using the following six training data samples.



x	y
5	64
6	56
12	50
9	71
15	44

$$\hat{y} = f(x) = w_0 + w_1x$$

A Regression Example

- Sum-of-squares error loss for training the linear model (finding the best w_0 and w_1):

$$\begin{aligned} O &= \frac{1}{2} \left[(\hat{y}_1 - y_1)^2 + (\hat{y}_2 - y_2)^2 + (\hat{y}_3 - y_3)^2 + (\hat{y}_4 - y_4)^2 + (\hat{y}_5 - y_5)^2 + (\hat{y}_6 - y_6)^2 \right] \\ &= \frac{1}{2} \sum_{i=1}^6 (\hat{y}_i - y_i)^2 \\ &= \frac{1}{2} \sum_{i=1}^6 (w_1 x_i + w_0 - y_i)^2 \end{aligned}$$

A Regression Example

- Incorporate the values of x_i and y_i to the error function:

$$\begin{aligned}x_1 &= 5, y_1 = 64 \\x_2 &= 6, y_2 = 56 \\x_3 &= 12, y_3 = 50 \\x_4 &= 9, y_4 = 71 \\x_5 &= 15, y_5 = 44 \\x_6 &= 16, y_6 = 60\end{aligned}$$


$$O(w_1, w_0) = \frac{1}{2} \left[(5w_1 + w_0 - 64)^2 + (6w_1 + w_0 - 56)^2 + (12w_1 + w_0 - 50)^2 + (9w_1 + w_0 - 71)^2 \right. \\ \left. + (15w_1 + w_0 - 44)^2 + (16w_1 + w_0 - 60)^2 \right]$$

- We want to minimise this training loss:

$$\min O(w_1, w_0)$$

A Regression Example

- Geometrically, to minimise this regression error loss enables you to find the best **red** line to have the shortest **blue** distances on average.



Linear Least Squares (LLS)

- **Linear least squares:** To train a linear model by minimising the sum-of-squares error.
 - Single-output case:

$$\min O(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N \left(y_i - \mathbf{w}^T \tilde{\mathbf{x}}_i \right)^2$$

- Multi-output case:

$$\min O(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^c \left(y_{ij} - \mathbf{w}_j^T \tilde{\mathbf{x}}_i \right)^2$$

Regularised Linear Least Squares

- A regularisation term can be added to the error function. For instance, in the single-output case, we have

$$\min O_{\lambda}(\mathbf{w}) = \text{sum of squares error} + \frac{\lambda}{2} \|\mathbf{w}\|_2^2 = \frac{1}{2} \sum_{i=1}^N (y_i - \mathbf{w}^T \tilde{\mathbf{x}}_i)^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

$$\mathbf{w}^T \mathbf{w} = \sum_{j=1}^{d+1} w_j^2$$

λ is a positive real-valued number set by the user.

- Other type of regulariser:

$$\min O_{\lambda}(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N (y_i - \mathbf{w}^T \tilde{\mathbf{x}}_i)^2 + \frac{\lambda}{2} \sum_{j=1}^{d+1} |w_j|^q$$

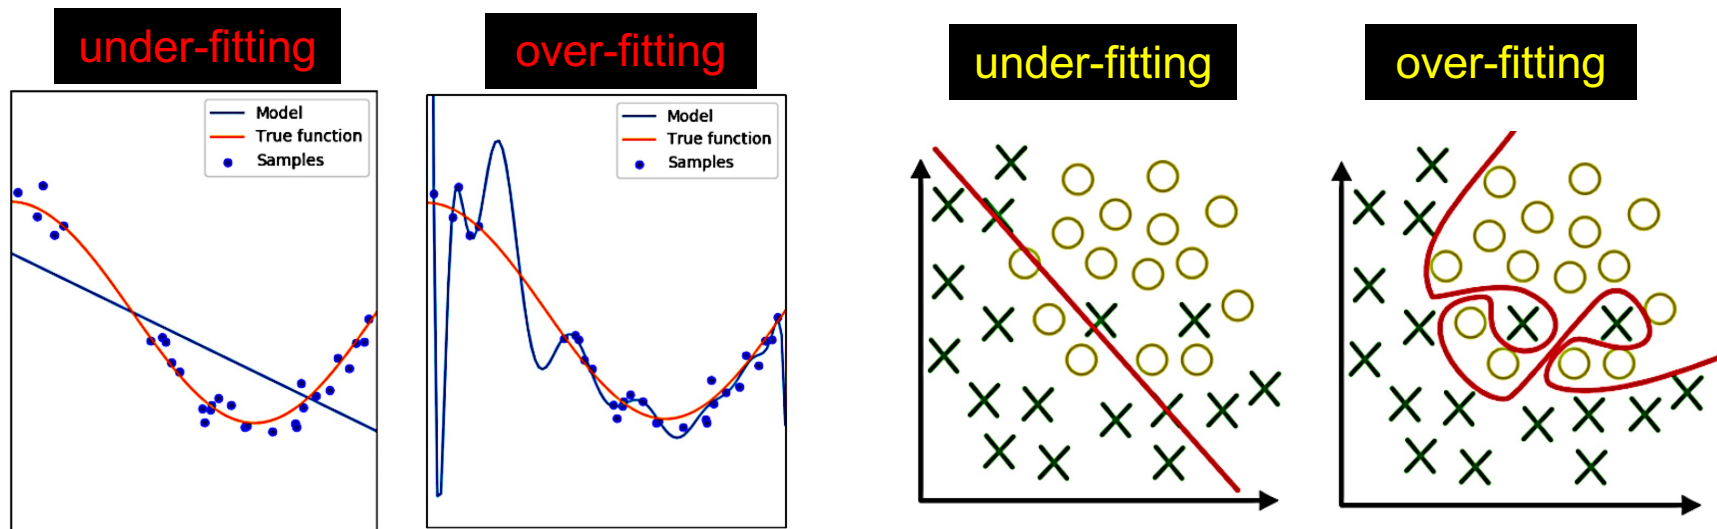
Here q is a positive integer set by the user.

- ❖ The case of $q=1$ (l_1 -regularisation) for regression is known as **lasso**.
- ❖ The case of $q=2$ (l_2 -regularisation) for regression is known as **ridge regression**.

- Regularisation prevents the model from over-fitting to training data.
- When λ is too large, it will lead to under-fitting though.

Why Regularisation?

- Over-fitting: Fit too closely to a particular set of data (e.g., training data), and may therefore fail to fit new data.
- Under-fitting: Cannot capture the underlying trend of the training data.
- Prevent over-fitting, e.g., $O_\lambda(\mathbf{w})$ gives less emphasis to the sum of squares error of the training data.



Probabilistic Regression Losses

Likelihood

- **Likelihood:** Given the observed data, it is the conditional probability assumed for the observed data given some parameter values.

$$\text{Likelihood}(\theta | \text{data}) = p(\text{data} | \theta)$$

- **Log likelihood:** Take the natural logarithm of the likelihood.

Likelihood Maximization

- A model can be trained by **maximising** the likelihood (or log likelihood) function of the training samples.
- Assume independence between samples.
 - **Maximum likelihood estimator (MLE):**

$$\max_{\theta} = \prod_{i=1}^N p(\mathbf{x}_i, y_i | \theta)$$

- **Log likelihood maximisation:**

$$\max_{\theta} = \sum_{i=1}^N \log p(\mathbf{x}_i, y_i | \theta)$$

Example: MLE for Linear Regression

- Likelihood of N training samples:

$$L = \prod_{i=1}^N N(y_i | \mathbf{w}^T \tilde{\mathbf{x}}_i, \sigma^2)$$

Assume a Gaussian distribution for likelihood estimation.

$$p(\text{data}_i | \theta) = p(y_i | \theta) = N(y_i | \mathbf{w}^T \tilde{\mathbf{x}}_i, \sigma^2)$$

- Log-likelihood:

$$O = \ln(L) = \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi) - \beta \frac{1}{2} \sum_{i=1}^N (y_i - \mathbf{w}^T \tilde{\mathbf{x}}_i)^2$$

where $\beta^{-1} = \sigma^2$

Sum-of-squares
error function!

- In this case, an MLE is equivalent to a sum-of-squares error minimizer.